

Reliable Conversational Analytics

Conversational Analytics becomes decision-grade when answer generation is guided by clear classification, disciplined narrative formulation, and deterministic validation.

★ Stakeholder trust is the product

Builds trust

Clear scope Honest refusal Correct numbers Traceable evidence Useful caveats

Breaks trust

Confident guesses Invented metrics Wrong query scope Narrative drift Hidden uncertainty

✓ How the architecture earns that trust

Before SQL

Check the question against a business-curated semantic model: valid use cases, metrics, field context, and known limits.

During SQL


Formulate safe queries against governed metadata and return structured evidence.

After SQL

Validate the narrative against the evidence: numbers, claims, caveats, and uncertainty.

→ Core workflow architecture

The inner layer is the agent trust core. The outer layer is the minimal production shell needed to run it responsibly.



∴ Summary

Reliable Conversational Analytics is not just text-to-SQL. It is a trust-controlled workflow that knows when to answer, when to clarify, and when to stop.

Project demo

I built an open-source Analytics Agent project demonstrating this pattern through a working app, GitHub repo, and LinkedIn write-up series.

[LinkedIn](#)

[GitHub repo](#)

[Live demo](#)